

The Photometric LSST Astronomical Time-series Classification Challenge (PLAsTiCC): Data set

Tarek Alam, Anita Bahmanyar, Rahul Biswas, Mi Dai, Lluís Galbany, Renée Hložek, Emille Ishida, Saurabh Jha, David Jones, Rick Kessler, Michelle Lochner, Ashish Mahabal, Kaisey Mandel, Juan Rafael Martinez Galarza, Alex Malz, Daniel Mutukrishna, Gautham Narayan, Tina Peters, Hiranya Peiris, and Kara Ponder

The Photometric LSST Astronomical Time Series Classification Challenge (PLAsTiCC) is an open data challenge to classify simulated astronomical time series data in preparation for the data from the Large Synoptic Survey Telescope (LSST), that will achieve first light in 2022. We briefly describe the PLAsTiCC data set that will be tested by the Kaggle team. This note will be updated for the full release of the data to the community.

1. Introduction

PLAsTiCC is a large data challenge where participants are asked to *classify astronomical time series data*. These simulated time series, or ‘light curves’ are measurements of an object’s brightness as a function of time - by counting the photon flux in six different astronomical filters. These filters include ultra-violet, optical and infrared regions of the wavelength spectrum. There are a large number of different types of astronomical objects, which make up different astronomical classes. The challenge is to analyse each set of light curves (1 light curve per filter, 6 filters per object) and classify each object as a member of such classes of objects. The time series data provided are simulations of what we expect from the upcoming Large Synoptic Survey Telescope (LSST), which will use an 8 meter telescope to image half the sky roughly once per week and over a ten year duration.

The users are asked to classify the data into 19 classes, 18 of which are represented in the training sample. The final class is meant to capture objects that are hypothesized to exist but have never been observed and are thus not in the training set.

In Figure 1, we show three example light curves from the training set. The top two panels show ‘transient’ objects which brighten and fade over a short time period. The lower panel shows a variable object which can fade temporarily, but always brightens again. Also note the gaps between observations: small gaps (days to weeks) from the time between telescope visits, and large gaps (> 6 months) where the object is not visible at night from the LSST site.

2. Astronomy Background

While we think of the night sky as static, it is filled with sources of light that vary in brightness on timescales from seconds and minutes to months and years.

Some of these events are classified as *transients*, and are the observational consequences of a large variety of astronomical phenomena. For example, the cataclysmic event that occurs when a star explodes generates a bright ‘supernova’ signal that fades with time, but does not repeat. Other events are classified as *variables*, since they vary repeatedly in brightness in a periodic (or aperiodic) fashion. Variable objects include emission from active galactic nuclei (AGN) at the hearts of galaxies, pulsating stars known as Cepheids, and eclipsing binary stars that alternate blocking out each other’s light from view.

These transient objects can provide important clues about themselves and their environment - as well as the evolution of the universe as a whole. For example, measurements of type Ia supernovae light curves provided the first evidence of accelerated expansion of the Universe, which might be caused by dark energy.

Each type of transient and variable provides a different clue that helps us study how stars evolve, the physics of stellar explosions, the chemical enrichment of the cosmos, and the accelerating expansion of the universe. Therefore, the proper classification of transients is a crucial task in observational astronomy - specially in light of the large data volumes expected for the next generation of astronomical surveys - which includes LSST.




Figure 1. Example light curves in the PLASTiCC data set. The three example objects display different changes in flux with time that are typical of real objects. The top-right panel illustrates that the brightening of the flux can occur near observation gaps, and therefore may not include the full time period of brightening (or fading) for the object. In addition, all three panels show that seasonal gaps and the instrument cadence of observations can introduce gaps in the light curve.

The question we address in this challenge is: *how well can we classify astronomical transients and variables from a light curve data set designed to mimic the data from LSST?* Crucially, the classification will occur on a large test set, but the training data will be a small, and poorly representative training set, to mimic the challenges we face observationally.

2.1. Different Methods for observing Astronomical Objects

Here we give more detail on LSST, and the challenge at hand. The two modes for characterising light from astronomical objects are called ‘spectroscopy’ and ‘photometry.’

Spectroscopy measures the flux per wavelength interval and is the modern equivalent of using a prism to separate a beam of light into a rainbow of colours. It is a high precision measurement which allows us to identify emission & absorption features indicative of specific chemical elements present in the object. Spectroscopy is also the most accurate and reliable tool that enables classification of astronomical transients and variables. Despite being paramount for the classification task, spectroscopy is an extremely time consuming process - with integration times ranging from 20 minutes to a few hours depending on the telescope and brightness of the source.

Given the volume of data expected from the upcoming large scale sky surveys, obtaining spectroscopy for every object is not feasible. An alternative approach is to take an image of the object through different ‘filters’, where each filter selects light of a specific colour, or wavelength range. For LSST there are six filters denoted u, g, r, i, z, Y , which select light from ultra-violet (u), green (g), red (r), and three near-infrared filters (izY). The filter efficiency vs. wavelength is shown in Fig. 2. For reference, the human eye is sensitive to light in the g & r bands. The flux of light in each filter, measured as a function of time, is a light curve. Classification is performed on these light curves. While a spectrum contains thousands of measurements in small wavelength regions, the light curve data includes at most six measurements (1 per filter) at any given time. The challenge is to classify objects with the highly compressed light curve information. Compared with spectroscopy, the advantage of measuring light curves is that we can observe objects that are much further away and much fainter. In addition, light curves from LSST can be measured over half the sky, a much larger region than spectroscopy can cover.

Beware that observations are sometimes degraded by moonlight, twilight, clouds, and wind. These degradations result in larger flux uncertainties, and this information is included in the light curve data (see Sec 3 below).

In addition to providing light curves, two other pieces of information are provided for each object. First is a proxy for distance called ‘redshift’, where more distant objects have a larger redshift. The training set includes accurate redshifts from the object, but the test data redshifts are approximate measurements based on *ugrizY* filter measurements from the ‘host galaxy’ of the object. While transients brighten and fade, the host galaxy fluxes don’t change and can thus be measured before an object’s light curve starts, or after it has faded. Beware that a few percent of the test data redshifts are catastrophic, meaning that some redshift uncertainties greatly under-estimate the difference between measured and true redshift.

The redshift effect is illustrated in Fig. 2. The black curve shows a *nearby* Type Ia supernova spectrum at a redshift of 0.01, corresponding to a distance of 140 million light years. While the term ‘nearby’ may seem strange in this case, this distance is indeed nearby when compared with the whole range of cosmic distances. Visual inspection of the supernova spectrum and the filter efficiencies shows that the maximum flux (spectrum summed over filter) is in the green (*g*) filter. The dashed curve shows a spectrum from a more distant Supernova, corresponding to a redshift of 0.5, or 5.1 billion light years away.¹ The maximum flux is now shifted to the red (*r*) filter. As the redshift and distance increase, the maximum flux appears in a redder filter: hence the term ‘redshift.’

The second piece of information is related to extinction from our Galaxy, known as the Milky Way. While our light curve measurements correct for the atmosphere and telescope transmission, we do not correct for absorption of light travelling through Milky Way ‘dust’ on its way to Earth. This absorption is strongest in the ultra-violet *u*-filter, and weakest in the infrared filters (*izY*). The data parameter has a strange name called ‘MWEBV’, which is an astronomical measure of how much redder an object appears compared to a Milky Way without dust. Larger MWEBV values correspond to more Milky Way dust, and objects appearing redder.

¹ The relation between distance and redshift is not linear, but is a function derived from General Relativity which depends on the properties of dark matter and dark energy.




Figure 2. Shaded regions show *ugrizY* filter efficiency vs. wavelength. Each measured filter flux is a sum of the photons collected within the wavelength range. Black curve shows a spectrum for a nearby Type Ia Supernova at redshift 0.01; dashed curve shows a spectrum for the same object at redshift 0.5. The dashed spectrum brightness is much lower than the solid spectrum, and has thus been scaled to see its shape relative to the solid spectrum.

3. The data

The light curve data for each object consists of a time series of fluxes in six filters (*ugrizY*), with non-uniform time sampling. The *PLAsTICC* data is provided in the form of a HDF5 file, which has two kinds of information. The first is a table listing each astronomical source in the data indexed by a unique identifier 'objid' which is a string. Each row of the table lists the properties of the source as follows:

- *objid*: the Object ID, unique identifier, string
- *ra*: right ascension, sky coordinate: co-longitude, units are degrees
- *dec1*: declination, sky coordinate: co-latitude, units are degrees

- `mwebv`: a property of the Milky Way dust along the line of sight to the astronomical source, and is thus a function of the sky coordinates of the source `ra`, `decl`. This is used to determine a `passband` dependent dimming and redenning of light from astronomical sources as described in subsection [2.1](#).
- `hostgal_specz`: the spectroscopic redshift of the source. This is an extremely accurate measure of redshift, and is not measured for the test sample. Hence, these values are null in the test data.
- `hostgal_photoz` : The photometric redshift of the host galaxy of the astronomical source. While this is meant to be a proxy for `hostgal_specz`, there can be large differences between the two and should be regarded as a far less accurate version of `hostgal_specz`.
- `hostgal_photoz_err` : The uncertainty on the `hostgal_photoz`
- `sntype` : The class of the astronomical source as a string. (currently this variable is called 'sntype' and will be changed in future)

The second table of information about the transients is its brightness as a function of time in different `passbands`, ie. light curve data. Each row of this table corresponds to an observation of the source at a particular time and `passband`. This table includes the following information

- `mjd`: the time (float) in Modified Julian Date (MJD) of the observation with a unit of day.
- `passband` : The specific LSST `passband` string: `u`, `g`, `r`, `i`, `z`, `Y` in which it was viewed.
- `flux`: the measured flux (brightness) in the `passband` of observation as listed in the `passband` column. This is a float.
- `fluxerr`: the uncertainty on the measurement of the `flux` listed above.
- `photflag`: ignore.

A few caveats about the light curve data are as follows:

- **Negative Flux** Due to statistical fluctuations and the way the brightness is estimated, the flux may be negative for dim sources, where the true flux is close to zero.

- **saturated observations** The light curves include ‘saturated’ observations of sources, where the source is too bright to obtain a measured value. In such cases, the `flux` is set to 0., and the `fluxerr` is set to 10,000,000. Such an observation may not yield a value of `flux`, but it indicates that the source was extremely bright rather than extremely dim at the time of observation.
- **Observing Cadences** The `objid` string has two prefixes DDF and ‘WFD. DDF corresponds to Deep Drill Fields over a small area of sky, but with very high quality light curves. WFD corresponds to wide-fast-deep over a very large sky area, but with lower quality light curves.

3.1. Obtaining the data and Scoring a classification

The header table is located in the sample HDF5 file at the path `header`, while the light curve is stored at the path `'objid'` in the sample HDF5 file. As part of the challenge, we provide an example Jupyter notebook to read in the data from an HDF5 file, and a notebook to compute the metrics for the challenge.

3.2. Training And Test Data

The training data follow the description above and have the properties and light curves of a set of 3200 astronomical sources and are meant to represent the brighter objects for which obtaining expensive spectroscopy is possible. The test data represent all of the data which have no spectroscopy. Therefore, the test data has 'NULL' entries for the two columns `hostgal_specz` and `sntype`. Moreover, for this reason their properties are **non-representative** of distributions of the the test data set. The training data are mostly comprised of nearby, low redshift, brighter samples while the test data contain higher redshift, fainter and more distant objects. Therefore, there are objects in the test data that do not have counterparts in the training data.

4. Challenge participation

PLAsTICC challenge entries are required to classify each of the sources in the header file of the test data set based on their properties and light curves. The classification must be done through the assignment of probabilities P_{ij} , the probability that the i th source in the test data is a member of the class j based on the combination of properties and light curve data for the object, the training data set, and any outside knowledge the participant may have acquired elsewhere. To specify the entry for a single source, the participant must provide the probabilities of that source belonging to each of the mutually exclusive (non-overlapping) 'classes in the training set, and of not belonging to any of the classes in the training set and therefore denoted by the `others` class. High values of P_{ij} for a particular class j and object i indicate that the participant believes that the i th source is likely to be a member of the j th class. As true probabilities, the quantities P_{ij} must satisfy

the following criteria:

$$0.0 \leq P_{ij} \leq 1.0 \quad \forall i, j$$

$$\sum_j P_{ij} = 1.0 \quad \forall i$$

For a PLAsTICC entry to be valid, it must have these probabilities for each class and astronomical source, ie. an entry cannot leave out probabilities on any source, or class. To win the challenge, the entry must be a valid entry and minimise the PLAsTiCC metric score (which is described in a separate note included in this challenge). For a participant using a classifier which decides that an object is of a particular class rather than provide probabilities, the participant has to use their own prescription to define probabilities. For example, a source classified as the first class may be given a probability of 1.0 and all other classes probabilities of 0.0, or the first class may be assigned 0.9 and all other classes may be given uniform probabilities so that the probabilities sum to unity.

For example, if the challenge was to classify a set of 3 observations into two classes of ‘star’ or ‘galaxy’ classes (and an ‘other’ class), the returned classification table would be 3×3 matrix:

Object ID	$P(star)$	$P(galaxy)$	$P(other)$
1	0.6	0.3	0.1
2	0.3	0.3	0.4
3	0.55	0.4	0.05

Table 1. An example classification table for a challenge to classify 3 objects into 3 classes

The PLAsTICC team involved in validating the data will not be able to participate in the challenge directly, and will only publish classifications on the data once the challenge has completed. Some PLAsTICC team members involved in defining metrics will participate in the challenge, but they have not seen any PLAsTICC information about the models or the data.

4.1. Acknowledgments

The PLAsTICC data relies on numerous members of the astronomical community to provide models of astronomical transients and variables. These models will be described in

a paper to be published once the challenge is complete. While we cannot thank them by name at this stage (as this could identify the models included in the challenge), we acknowledge their contributions anonymously at this stage. This work was supported by an LSST Corporation Enabling Science grant, and a Dark Energy Science Collaboration Workshop support grant.

The DESC acknowledges ongoing support from the Institut National de Physique Nucléaire et de Physique des Particules in France; the Science & Technology Facilities Council in the United Kingdom; and the Department of Energy, the National Science Foundation, and the LSST Corporation in the United States. DESC uses resources of the IN2P3 Computing Center (CC-IN2P3—Lyon/Villeurbanne - France) funded by the Centre National de la Recherche Scientifique; the National Energy Research Scientific Computing Center, a DOE Office of Science User Facility supported by the Office of Science of the U.S. Department of Energy under Contract No. DE-AC02-05CH11231; STFC DiRAC HPC Facilities, funded by UK BIS National E-infrastructure capital grants; and the UK particle physics grid, supported by the GridPP Collaboration. This work was performed in part under DOE Contract DE-AC02-76SF00515.